

# The Portable Extensible Toolkit for Scientific Computing

Matthew Knepley

Mathematics and Computer Science Division  
Argonne National Laboratory

Computation Institute  
University of Chicago

ACTS Workshop  
NERSC, Berkeley, CA      August 18, 2010



# Outline

- 1 Getting Started with PETSc
  - What is PETSc?
  - Who uses PETSc?
  - How can I get PETSc?
  - How do I Configure PETSc?
  - How do I Build PETSc?
  - How do I run an example?
  - How do I get more help?

2 PETSc Integration

3 Common PETSc Usage

4 Advanced PETSc

# Outline

## 1 Getting Started with PETSc

- What is PETSc?
- Who uses PETSc?
- How can I get PETSc?
- How do I Configure PETSc?
- How do I Build PETSc?
- How do I run an example?
- How do I get more help?

# Unit Objectives

- Introduce PETSc
- Download, Configure, Build, and Run an Example
- Empower students to learn more about PETSc

# What I Need From You

- Tell me if you do not understand
- Tell me if an example does not work
- Suggest better wording or **figures**
- Followup problems at [petsc-maint@mcs.anl.gov](mailto:petsc-maint@mcs.anl.gov)

# How We Can Help at the Tutorial

- Point out relevant documentation
- Quickly answer questions
- Help install
- Guide design of large scale codes
- Answer email at [petsc-maint@mcs.anl.gov](mailto:petsc-maint@mcs.anl.gov)

# How We Can Help at the Tutorial

- Point out relevant documentation
- Quickly answer questions
- Help install
- Guide design of large scale codes
- Answer email at [petsc-maint@mcs.anl.gov](mailto:petsc-maint@mcs.anl.gov)

# How We Can Help at the Tutorial

- Point out relevant documentation
- Quickly answer questions
- Help install
- Guide design of large scale codes
- Answer email at [petsc-maint@mcs.anl.gov](mailto:petsc-maint@mcs.anl.gov)



# How We Can Help at the Tutorial

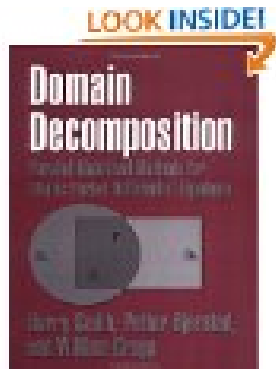
- Point out relevant documentation
- Quickly answer questions
- Help install
- Guide design of large scale codes
- Answer email at [petsc-maint@mcs.anl.gov](mailto:petsc-maint@mcs.anl.gov)

# How did PETSc Originate?

PETSc was developed as a Platform for  
**Experimentation**

We want to experiment with different

- Models
- Discretizations
- Solvers
- Algorithms
  - which blur these boundaries



# The Role of PETSc

*Developing parallel, nontrivial PDE solvers that deliver high performance is still difficult and requires months (or even years) of concentrated effort.*

*PETSc is a toolkit that can ease these difficulties and reduce the development time, but it is not a black-box PDE solver, nor a **silver bullet**.*

— Barry Smith

# What is PETSc?

A freely available and supported research code

- Download from <http://www.mcs.anl.gov/petsc>
- Free for everyone, including industrial users
- Hyperlinked manual, examples, and manual pages for all routines
- Hundreds of tutorial-style examples
- Support via email: [petsc-maint@mcs.anl.gov](mailto:petsc-maint@mcs.anl.gov)
- Usable from C, C++, Fortran 77/90, and Python

# What is PETSc?

- Portable to any parallel system supporting MPI, including:
  - Tightly coupled systems
    - Cray XT5, BG/P, NVIDIA Tesla, Earth Simulator, Sun Blade
  - Loosely coupled systems, such as networks of workstations
    - IBM, Mac, Sun, PCs running Linux or Windows
- PETSc History
  - Begun September 1991
  - Over 60,000 downloads since 1995 (version 2)
  - Currently 400 per month
- PETSc Funding and Support
  - Department of Energy
    - SciDAC, MICS Program, INL Reactor Program
  - National Science Foundation
    - CIG, CISE, Multidisciplinary Challenge Program

# The PETSc Team



Bill Gropp



Barry Smith



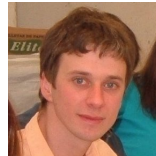
Satish Balay



Jed Brown



Matt Knepley



Lisandro Dalcin



Hong Zhang



Victor Eijkhout



Dmitry Karpeev

# Outline

## 1 Getting Started with PETSc

- What is PETSc?
- **Who uses PETSc?**
- How can I get PETSc?
- How do I Configure PETSc?
- How do I Build PETSc?
- How do I run an example?
- How do I get more help?

# Who Uses PETSc?

- Computational Scientists

- PyLith (CIG), Underworld (Monash), Magma Dynamics (LDEO, Columbia), PFLOTRAN (DOE)

- Algorithm Developers

- Iterative methods and Preconditioning researchers

- Package Developers

- SLEPc, TAO, DealII, PETSc-FEM, MagPar, PetFMM, PetRBF



# What Can We Handle?

- PETSc has run implicit problems with over **1 billion** unknowns
  - PFLOTRAN for flow in porous media
- PETSc has run on over **224,000** cores efficiently
  - UNIC on the IBM BG/P Intrepid at ANL
  - PFLOTRAN on the Cray XT5 Jaguar at ORNL
- PETSc applications have run at **3 Teraflops**
  - LANL PFLOTRAN code

# What Can We Handle?

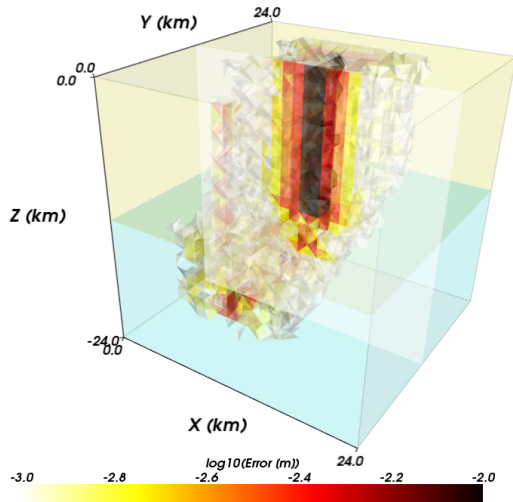
- PETSc has run implicit problems with over **1 billion** unknowns
  - PFLOTRAN for flow in porous media
- PETSc has run on over **224,000** cores efficiently
  - UNIC on the IBM BG/P Intrepid at ANL
  - PFLOTRAN on the Cray XT5 Jaguar at ORNL
- PETSc applications have run at **3 Teraflops**
  - LANL PFLOTRAN code

# What Can We Handle?

- PETSc has run implicit problems with over **1 billion** unknowns
  - PFLOTRAN for flow in porous media
- PETSc has run on over **224,000** cores efficiently
  - UNIC on the IBM BG/P Intrepid at ANL
  - PFLOTRAN on the Cray XT5 Jaguar at ORNL
- PETSc applications have run at **3 Teraflops**
  - LANL PFLOTRAN code

# PyLith

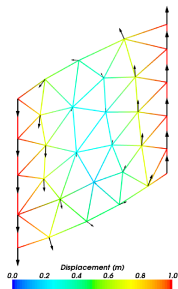
- Multiple problems
  - Dynamic rupture
  - Quasi-static relaxation
- Multiple models
  - Nonlinear visco-plastic
  - Finite deformation
  - Fault constitutive models
- Multiple meshes
  - 1D, 2D, 3D
  - Hex and tet meshes
- Parallel
  - PETSc solvers
  - Sieve mesh management



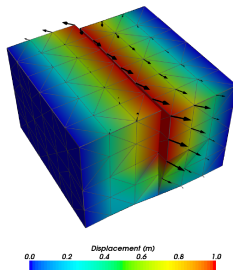
<sup>a</sup>Aagaard, Knepley, Williams

# Multiple Mesh Types

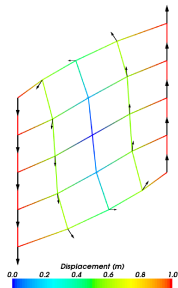
Triangular



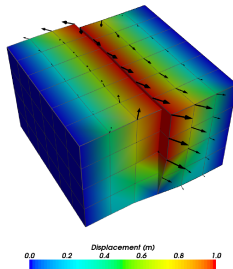
Tetrahedral



Rectangular

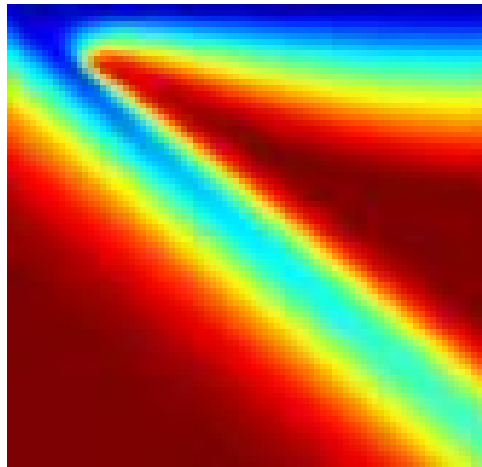


Hexahedral



# Magma Dynamics

- Couples scales
  - Subduction
  - Magma Migration
- Physics
  - Incompressible fluid
  - Porous solid
  - Variable porosity
- Deforming matrix
  - Compaction pressure
- Code generation
  - FEniCS
- Multiphysics Preconditioning
  - PETSc FieldSplit

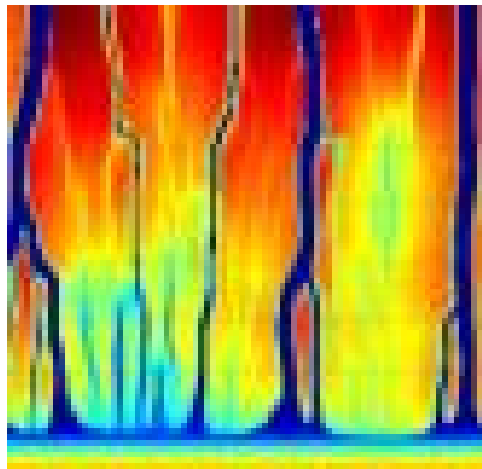


---

<sup>a</sup>Katz, Spiegelman

# Magma Dynamics

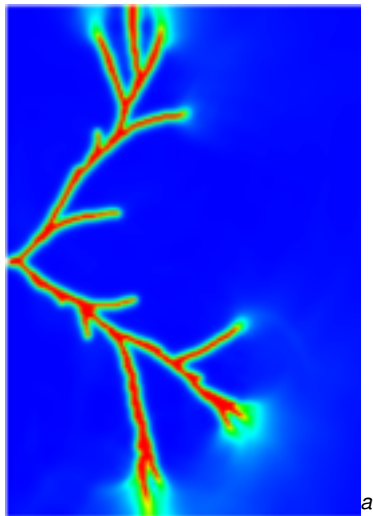
- Couples scales
  - Subduction
  - Magma Migration
- Physics
  - Incompressible fluid
  - Porous solid
  - Variable porosity
- Deforming matrix
  - Compaction pressure
- Code generation
  - FEniCS
- Multiphysics Preconditioning
  - PETSc FieldSplit



<sup>a</sup>Katz, Spiegelman

# Fracture Mechanics

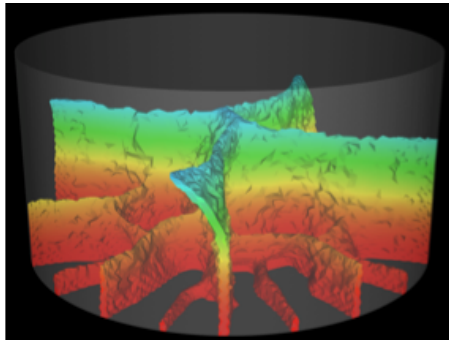
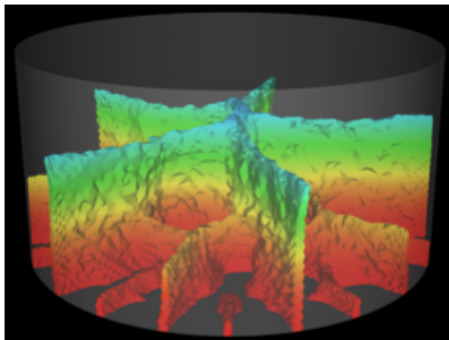
- Full variational formulation
  - Phase field
  - Linear or Quadratic penalty
- Uses TAO optimization
  - Necessary for linear penalty
  - Backtracking
- No prescribed cracks
  - Arbitrary crack geometry
  - Arbitrary intersections
- Multiple materials
  - Composite toughness



<sup>a</sup>Bourdin



# Fracture Mechanics



1

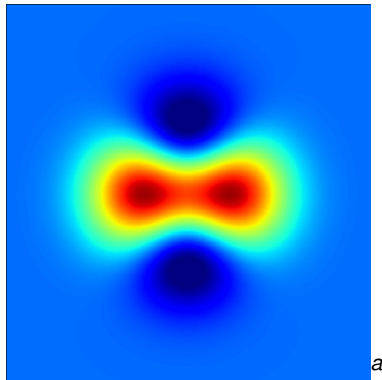
---

<sup>1</sup>Bourdin

# Vortex Method

$t = 000$

- Incompressible Flow
  - Gaussian vortex blobs
  - High Re
- PetFMM
  - 2D/3D domains
  - Automatic load balancing
  - Variety of kernels
  - Optimized with templates
- PetRBF
  - Variety of RBFs
  - Uses PETSc solvers
  - Scalable preconditioner
- Parallelism
  - MPI
  - GPU

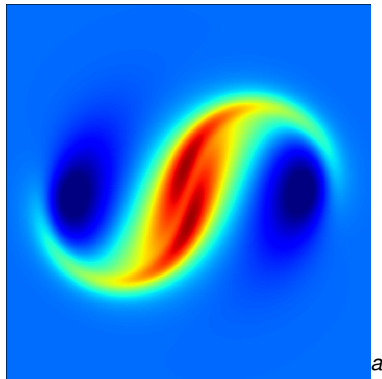


<sup>a</sup>Cruz, Yokota, Barba, Knepley

# Vortex Method

$t = 100$

- Incompressible Flow
  - Gaussian vortex blobs
  - High Re
- PetFMM
  - 2D/3D domains
  - Automatic load balancing
  - Variety of kernels
  - Optimized with templates
- PetRBF
  - Variety of RBFs
  - Uses PETSc solvers
  - Scalable preconditioner
- Parallelism
  - MPI
  - GPU

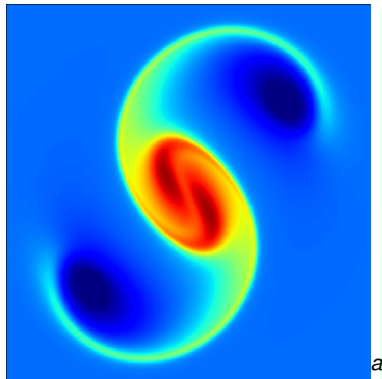


<sup>a</sup>Cruz, Yokota, Barba, Knepley

# Vortex Method

$t = 200$

- Incompressible Flow
  - Gaussian vortex blobs
  - High Re
- PetFMM
  - 2D/3D domains
  - Automatic load balancing
  - Variety of kernels
  - Optimized with templates
- PetRBF
  - Variety of RBFs
  - Uses PETSc solvers
  - Scalable preconditioner
- Parallelism
  - MPI
  - GPU

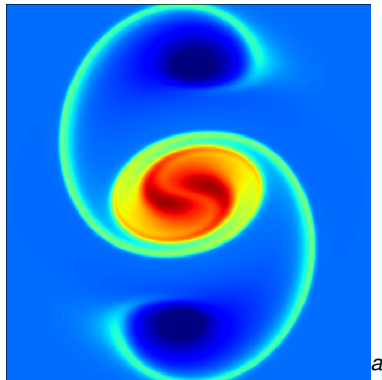


<sup>a</sup>Cruz, Yokota, Barba, Knepley

# Vortex Method

$t = 300$

- Incompressible Flow
  - Gaussian vortex blobs
  - High Re
- PetFMM
  - 2D/3D domains
  - Automatic load balancing
  - Variety of kernels
  - Optimized with templates
- PetRBF
  - Variety of RBFs
  - Uses PETSc solvers
  - Scalable preconditioner
- Parallelism
  - MPI
  - GPU

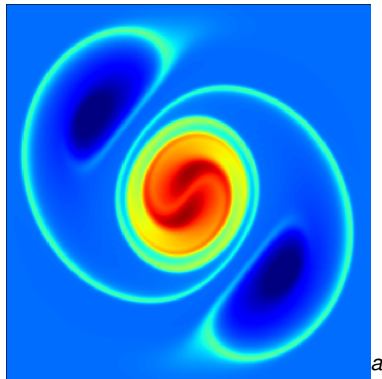


<sup>a</sup>Cruz, Yokota, Barba, Knepley

# Vortex Method

$t = 400$

- Incompressible Flow
  - Gaussian vortex blobs
  - High Re
- PetFMM
  - 2D/3D domains
  - Automatic load balancing
  - Variety of kernels
  - Optimized with templates
- PetRBF
  - Variety of RBFs
  - Uses PETSc solvers
  - Scalable preconditioner
- Parallelism
  - MPI
  - GPU

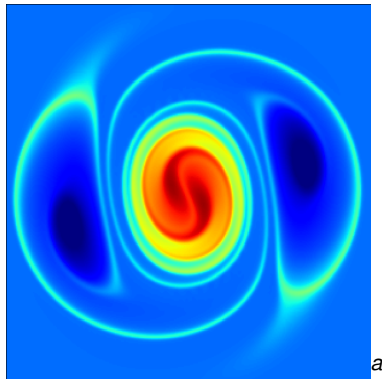


<sup>a</sup>Cruz, Yokota, Barba, Knepley

# Vortex Method

$t = 500$

- Incompressible Flow
  - Gaussian vortex blobs
  - High Re
- PetFMM
  - 2D/3D domains
  - Automatic load balancing
  - Variety of kernels
  - Optimized with templates
- PetRBF
  - Variety of RBFs
  - Uses PETSc solvers
  - Scalable preconditioner
- Parallelism
  - MPI
  - GPU

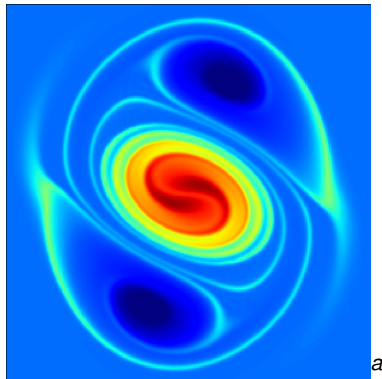


<sup>a</sup>Cruz, Yokota, Barba, Knepley

# Vortex Method

$t = 600$

- Incompressible Flow
  - Gaussian vortex blobs
  - High Re
- PetFMM
  - 2D/3D domains
  - Automatic load balancing
  - Variety of kernels
  - Optimized with templates
- PetRBF
  - Variety of RBFs
  - Uses PETSc solvers
  - Scalable preconditioner
- Parallelism
  - MPI
  - GPU



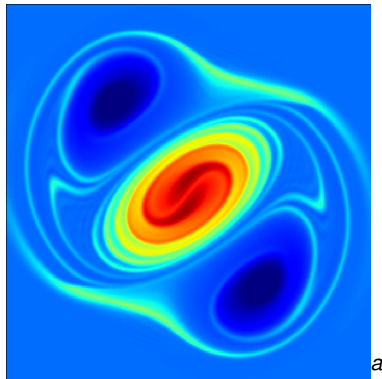
<sup>a</sup>Cruz, Yokota, Barba, Knepley



# Vortex Method

$t = 700$

- Incompressible Flow
  - Gaussian vortex blobs
  - High Re
- PetFMM
  - 2D/3D domains
  - Automatic load balancing
  - Variety of kernels
  - Optimized with templates
- PetRBF
  - Variety of RBFs
  - Uses PETSc solvers
  - Scalable preconditioner
- Parallelism
  - MPI
  - GPU

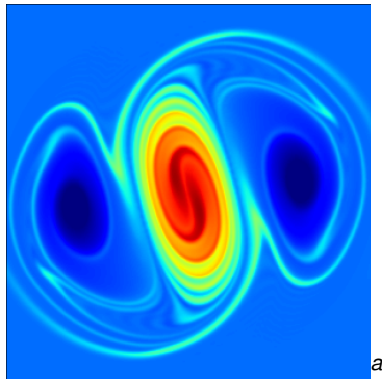


<sup>a</sup>Cruz, Yokota, Barba, Knepley

# Vortex Method

$t = 800$

- Incompressible Flow
  - Gaussian vortex blobs
  - High Re
- PetFMM
  - 2D/3D domains
  - Automatic load balancing
  - Variety of kernels
  - Optimized with templates
- PetRBF
  - Variety of RBFs
  - Uses PETSc solvers
  - Scalable preconditioner
- Parallelism
  - MPI
  - GPU

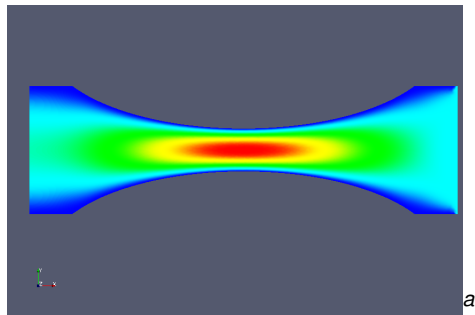


<sup>a</sup>Cruz, Yokota, Barba, Knepley

# FEniCS-Apps

Dolfin-grade2

- Rheologies
  - Maxwell
  - Grade 2
  - Oldroyd-B
- Stabilization
  - DG
  - SUPG
  - EVSS
  - DEVSS
  - Macroelement
- Automation
  - FIAT (elements)
  - FFC (weak forms)

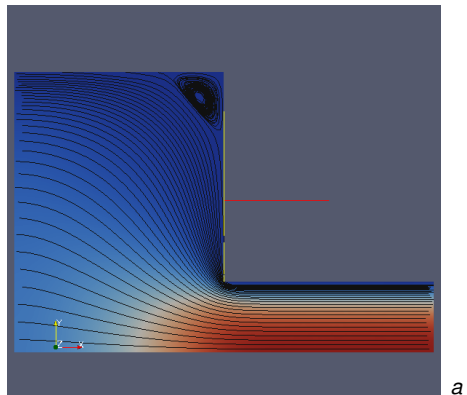


<sup>a</sup>Terrel

# FEniCS-Apps

## Dolfin-grade2

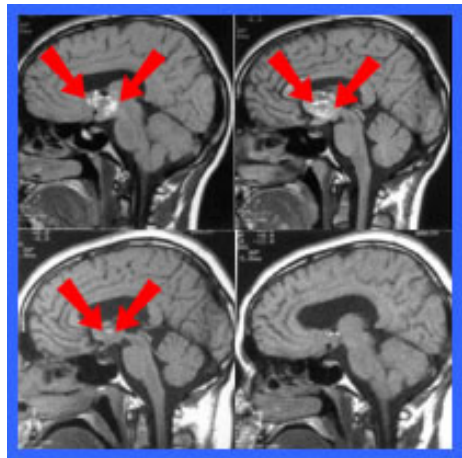
- Rheologies
  - Maxwell
  - Grade 2
  - Oldroyd-B
- Stabilization
  - DG
  - SUPG
  - EVSS
  - DEVSS
  - Macroelement
- Automation
  - FIAT (elements)
  - FFC (weak forms)



<sup>a</sup>Terrel

# Real-time Surgery

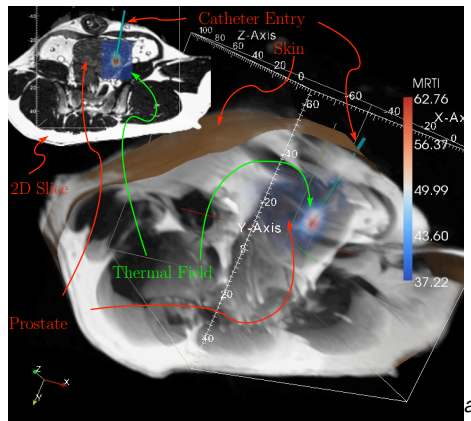
- Brain Surgery
  - Elastic deformation
  - Overlaid on MRI
  - Guides surgeon
- Laser Thermal Therapy
  - PDE constrained optimization
  - Per-patient calibration
  - Thermal inverse problem



<sup>a</sup>Warfield, Ferrant, et.al.

# Real-time Surgery

- Brain Surgery
  - Elastic deformation
  - Overlaid on MRI
  - Guides surgeon
- Laser Thermal Therapy
  - PDE constrained optimization
  - Per-patient calibration
  - Thermal inverse problem



<sup>a</sup>Fuentes, Oden, et.al.

# Outline

## 1 Getting Started with PETSc

- What is PETSc?
- Who uses PETSc?
- **How can I get PETSc?**
- How do I Configure PETSc?
- How do I Build PETSc?
- How do I run an example?
- How do I get more help?

# Downloading PETSc

- The latest tarball is on the PETSc site
  - <ftp://ftp.mcs.anl.gov/pub/petsc/petsc.tar.gz>
  - We no longer distribute patches (everything is in the distribution)
- There is a Debian package
- There is a FreeBSD Port
- There is a Mercurial development repository



# Cloning PETSc

- The full development repository is open to the public
  - <http://petsc.cs.iit.edu/petsc/petsc-dev>
  - <http://petsc.cs.iit.edu/petsc/BuildSystem>
- Why is this better?
  - You can clone to any release (or any specific ChangeSet)
  - You can easily rollback changes (or releases)
  - You can get fixes from us the same day
- We also make release repositories available
  - <http://petsc.cs.iit.edu/petsc/releases/petsc-3.1>
  - <http://petsc.cs.iit.edu/petsc/releases/BuildSystem-3.1>

# Unpacking PETSc

- Just clone development repository

- `hg clone http://petsc.cs.iit.edu/petsc/petsc-dev  
petsc-dev`
- `hg clone -rrelease-3.1 petsc-dev petsc-3.1`

**or**

- Unpack the tarball

- `tar xzf petsc.tar.gz`

# Exercise 1

Download and Unpack PETSc!

# Outline

## 1 Getting Started with PETSc

- What is PETSc?
- Who uses PETSc?
- How can I get PETSc?
- **How do I Configure PETSc?**
- How do I Build PETSc?
- How do I run an example?
- How do I get more help?

# Configuring PETSc

- Set `$PETSC_DIR` to the installation root directory
- Run the configuration utility
  - `$PETSC_DIR/configure`
  - `$PETSC_DIR/configure -help`
  - `$PETSC_DIR/configure -download-mpich`
  - `$PETSC_DIR/configure -prefix=/usr`
- There are many examples on the installation page
- Configuration files are in `$PETSC_DIR/$PETSC_ARCH/conf`
  - Configure header is in `$PETSC_DIR/$PETSC_ARCH/include`
  - `$PETSC_ARCH` has a default if not specified

# Configuring PETSc

- You can easily reconfigure with the same options
  - `./$PETSC_ARCH/conf/reconfigure-$PETSC_ARCH.py`
- Can maintain several different configurations
  - `./configure -PETSC_ARCH=linux-fast`  
`-with-debugging=0`
- All configuration information is in the logfile
  - `./$PETSC_ARCH/conf/configure.log`
  - **ALWAYS** send this file with bug reports

# Automatic Downloads

- Starting in 2.2.1, some packages are automatically
  - Downloaded
  - Configured and Built (in `$PETSC_DIR/externalpackages`)
  - Installed with PETSc
- Currently works for
  - petsc4py
  - PETSc documentation utilities (Sowing, lgrind, c2html)
  - BLAS, LAPACK, BLACS, ScaLAPACK, PLAPACK
  - MPICH, MPE, LAM
  - ParMetis, Chaco, Jostle, Party, Scotch, Zoltan
  - MUMPS, Spooles, SuperLU, SuperLU\_Dist, UMFPack, pARMS
  - BLOPEX, FFTW, SPRNG
  - Prometheus, HYPRE, ML, SPAI
  - Sundials
  - Triangle, TetGen
  - FIAT, FFC, Generator
  - Boost

# Exercise 2

Configure your downloaded PETSc.



# Outline

## 1 Getting Started with PETSc

- What is PETSc?
- Who uses PETSc?
- How can I get PETSc?
- How do I Configure PETSc?
- **How do I Build PETSc?**
- How do I run an example?
- How do I get more help?

# Building PETSc

- Uses recursive make starting in `cd $PETSC_DIR`
  - `make`
  - `make install` if you configured with `--prefix`
  - Check build when done with `make test`
- Complete log for each build is in logfile
  - `./$PETSC_ARCH/conf/make.log`
  - ALWAYS send this with bug reports
- Can build multiple configurations
  - `PETSC_ARCH=linux-fast make`
  - Libraries are in `$PETSC_DIR/$PETSC_ARCH/lib/`
- Can also build a subtree
  - `cd src/snes; make`
  - `cd src/snes; make ACTION=libfast tree`

# Exercise 3

Build your configured PETSc.

# Exercise 4

## Reconfigure PETSc to use ParMetis.

1

```
linux-gnu-c-debug/conf/reconfigure-linux-gnu-c-debug.py
```

- `-PETSC_ARCH=linux-parmetis`
- `-download-parmetis`

2

```
PETSC_ARCH=linux-parmetis make
```

3

```
PETSC_ARCH=linux-parmetis make test
```

# Outline

## 1 Getting Started with PETSc

- What is PETSc?
- Who uses PETSc?
- How can I get PETSc?
- How do I Configure PETSc?
- How do I Build PETSc?
- **How do I run an example?**
- How do I get more help?

# Running PETSc

- Try running PETSc examples first
  - `cd $PETSC_DIR/src/snes/examples/tutorials`
- Build examples using make targets
  - `make ex5`
- Run examples using the make target
  - `make runex5`
- Can also run using MPI directly
  - `mpirun ./ex5 -snes_max_it 5`
  - `mpiexec ./ex5 -snes_monitor`

# Using MPI

- The **M**essage **P**assing **I**nterface is:
  - a library for parallel communication
  - a system for launching parallel jobs (mpirun/mpiexec)
  - a community standard
- Launching jobs is easy
  - `mpiexec -n 4 ./ex5`
- You should never have to make MPI calls when using PETSc
  - Almost never

# MPI Concepts

- Communicator
  - A context (or scope) for parallel communication (“Who can I talk to”)
  - There are two defaults:
    - yourself (PETSC\_COMM\_SELF),
    - and everyone launched (PETSC\_COMM\_WORLD)
  - Can create new communicators by splitting existing ones
  - Every PETSc object has a communicator
  - Set PETSC\_COMM\_WORLD to put all of PETSc in a subcomm
- Point-to-point communication
  - Happens between two processes (like in `MatMult()`)
- Reduction or scan operations
  - Happens among all processes (like in `VecDot()`)



# Alternative Memory Models

- Single process (address space) model
  - OpenMP and threads in general
  - Fortran 90/95 and compiler-discovered parallelism
  - System manages memory and (usually) thread scheduling
  - Named variables refer to the same storage
- Single name space model
  - HPF, UPC
  - Global Arrays
  - Titanium
  - Variables refer to the coherent values (distribution is automatic)
- Distributed memory (shared nothing)
  - Message passing
  - Names variables in different processes are unrelated

# Common Viewing Options

- Gives a text representation
  - `-vec_view`
- Generally views subobjects too
  - `-snes_view`
- Can visualize some objects
  - `-mat_view_draw`
- Alternative formats
  - `-vec_view_binary`, `-vec_view_matlab`,  
`-vec_view_socket`
- Sometimes provides extra information
  - `-mat_view_info`, `-mat_view_info_detailed`

# Common Monitoring Options

- Display the residual
  - `-ksp_monitor`, graphically `-ksp_monitor_draw`
- Can disable dynamically
  - `-ksp_monitors_cancel`
- Does not display subsolvers
  - `-snes_monitor`
- Can use the true residual
  - `-ksp_monitor_true_residual`
- Can display different subobjects
  - `-snes_monitor_residual`, `-snes_monitor_solution`,  
`-snes_monitor_solution_update`
  - `-snes_monitor_range`
  - `-ksp_gmres_krylov_monitor`
- Can display the spectrum
  - `-ksp_monitor_singular_value`

# Exercise 5

Run SNES Example 5 using some custom options.

- ❶ `cd $PETSC_DIR/src/snes/examples/tutorials`
- ❷ `make ex5`
- ❸ `mpiexec ./ex5 -snes_monitor -snes_view`
- ❹ `mpiexec ./ex5 -snes_type tr -snes_monitor  
-snes_view`
- ❺ `mpiexec ./ex5 -ksp_monitor -snes_monitor  
-snes_view`
- ❻ `mpiexec ./ex5 -pc_type jacobi -ksp_monitor  
-snes_monitor -snes_view`
- ❼ `mpiexec ./ex5 -ksp_type bicg -ksp_monitor  
-snes_monitor -snes_view`

# Exercise 6

## Create a new code based upon SNES Example 5.

### 1 Create a new directory

- `mkdir -p /home/knepley/proj/newsim/src`

### 2 Copy the source

- `cp ex5.c /home/knepley/proj/newsim/src`
- **Add** `myStuff.c` and `myStuff2.F`

### 3 Create a PETSc makefile

- `bin/ex5: src/ex5.o src/myStuff.o src/myStuff2.o`
- `{CLINKER} -o $@ $^ ${PETSC_SNES_LIB}`
- `include ${PETSC_DIR}/conf/variables`
- `include ${PETSC_DIR}/conf/rules`

To get the project ready-made

```
hg clone
```

```
http://petsc.cs.iit.edu/petsc/tutorials/SimpleTutorial  
newsim
```

# Outline

## 1 Getting Started with PETSc

- What is PETSc?
- Who uses PETSc?
- How can I get PETSc?
- How do I Configure PETSc?
- How do I Build PETSc?
- How do I run an example?
- How do I get more help?

# Getting More Help

- <http://www.mcs.anl.gov/petsc>
- Hyperlinked documentation
  - Manual
  - Manual pages for every method
  - HTML of all example code (linked to manual pages)
- FAQ
- Full support at [petsc-maint@mcs.anl.gov](mailto:petsc-maint@mcs.anl.gov)
- High profile users
  - David Keyes
  - Marc Spiegelman
  - Richard Katz
  - Brad Aagaard
  - Lorena Barba
  - Jed Brown

# Outline

## 1 Getting Started with PETSc

## 2 PETSc Integration

- Initial Operations
- Vector Algebra
- Matrix Algebra
- Algebraic Solvers
- More Abstractions

## 3 Common PETSc Usage

## 4 Advanced PETSc

## 5 Future Plans



# Outline

2

## PETSc Integration

- Initial Operations
- Vector Algebra
- Matrix Algebra
- Algebraic Solvers
- More Abstractions

# Application Integration

- Be willing to experiment with algorithms
  - No optimality without interplay between physics and algorithmics
- Adopt flexible, extensible programming
  - Algorithms and data structures not hardwired
- Be willing to play with the real code
  - Toy models are rarely helpful
- If possible, profile before integration
  - Automatic in PETSc

# PETSc Integration

PETSc is a set a library interfaces

- We do not seize `main()`
- We do not control output
- We propagate errors from underlying packages
- We present the same interfaces in:
  - C
  - C++
  - F77
  - F90
  - Python

See Gropp in [SIAM, OO Methods for Interop SciEng, '99](#)

# Integration Stages

- **Version Control**
  - It is impossible to overemphasize
  - We use **Mercurial**
- Initialization
  - Linking to PETSc
- Profiling
  - Profile **before** changing
  - Also incorporate command line processing
- Linear Algebra
  - First PETSc data structures
- Solvers
  - Very easy after linear algebra is integrated

# Initialization

- Call `PetscInitialize()`
  - Setup static data and services
  - Setup MPI if it is not already
- Call `PetscFinalize()`
  - Calculates logging summary
  - Shutdown and release resources
- Checks compile and link

# Profiling

- Use `-log_summary` for a performance profile
  - Event timing
  - Event flops
  - Memory usage
  - MPI messages
- Call `PetscLogStagePush()` and `PetscLogStagePop()`
  - User can add new stages
- Call `PetscLogEventBegin()` and `PetscLogEventEnd()`
  - User can add new events

# Command Line Processing

- Check for an option
  - `PetscOptionsHasName()`
- Retrieve a value
  - `PetscOptionsGetInt()`, `PetscOptionsGetIntArray()`
- Set a value
  - `PetscOptionsSetValue()`
- Check for unused options
  - `-options_left`
- Clear, alias, reject, etc.
- Modern form uses
  - `PetscOptionsBegin()`, `PetscOptionsEnd()`
  - `PetscOptionsInt()`, `PetscOptionsReal()`
  - Integrates with `-help`

# Outline

2

## PETSc Integration

- Initial Operations
- **Vector Algebra**
- Matrix Algebra
- Algebraic Solvers
- More Abstractions



# Vector Algebra

## What are PETSc vectors?

- Fundamental objects representing field solutions, right-hand sides, etc.
- Each process locally owns a subvector of contiguous global data

## How do I create vectors?

- `VecCreate (MPI_Comm, Vec *)`
- `VecSetSizes (Vec, int n, int N)`
- `VecSetType (Vec, VecType typeName)`
- `VecSetFromOptions (Vec)`
  - Can set the type at runtime

# Vector Algebra

## A PETSc Vec

- Has a direct interface to the values
- Supports all vector space operations
  - `VecDot()`, `VecNorm()`, `VecScale()`
- Has unusual operations, e.g.
  - `VecSqrt()`, `VecWhichBetween()`
- Communicates automatically during assembly
- Has customizable communication (scatters)

# Parallel Assembly

## Vectors and Matrices

- Processes may set an arbitrary entry
  - Must use proper interface
- Entries need not be generated locally
  - Local meaning the process on which they are stored
- PETSc automatically moves data if necessary
  - Happens during the assembly phase

# Vector Assembly

- A three step process
  - Each process sets or adds values
  - Begin communication to send values to the correct process
  - Complete the communication
- `VecSetValues(Vec v, int n, int rows[], PetscScalar values[], mode)`
  - mode is either  
`INSERT_VALUES`, `ADD_VALUES`
- Two phase assembly allows overlap of communication and computation
  - `VecAssemblyBegin(Vec v)`
  - `VecAssemblyEnd(Vec v)`

# One Way to Set the Elements of a Vector

```
VecGetSize(x, &N);  
MPI_Comm_rank(PETSC_COMM_WORLD, &rank);  
if (rank == 0) {  
    val = 0.0;  
    for(i = 0; i < N; ++i) {  
        VecSetValues(x, 1, &i, &val, INSERT_VALUES);  
        val += 10.0;  
    }  
}  
  
/* These routines ensure that the data is  
   distributed to the other processes */  
VecAssemblyBegin(x);  
VecAssemblyEnd(x);
```

# A Better Way to Set the Elements of a Vector

```
VecGetOwnershipRange(x, &low, &high);  
val = low*10.0;  
for(i = low; i < high; ++i) {  
    VecSetValues(x, 1, &i, &val, INSERT_VALUES);  
    val += 10.0;  
}  
  
/* These routines ensure that the data is  
   distributed to the other processes */  
VecAssemblyBegin(x);  
VecAssemblyEnd(x);
```

# Selected Vector Operations

Function Name	Operation
VecAXPY(Vec y, PetscScalar a, Vec x)	$y = y + a * x$
VecAYPX(Vec y, PetscScalar a, Vec x)	$y = x + a * y$
VecWAYPX(Vec w, PetscScalar a, Vec x, Vec y)	$w = y + a * x$
VecScale(Vec x, PetscScalar a)	$x = a * x$
VecCopy(Vec y, Vec x)	$y = x$
VecPointwiseMult(Vec w, Vec x, Vec y)	$w_i = x_i * y_i$
VecMax(Vec x, PetscInt *idx, PetscScalar *r)	$r = \max r_i$
VecShift(Vec x, PetscScalar r)	$x_i = x_i + r$
VecAbs(Vec x)	$x_i =  x_i $
VecNorm(Vec x, NormType type, PetscReal *r)	$r =   x  $

# Working With Local Vectors

It is sometimes more efficient to directly access local storage of a `Vec`.

- PETSc allows you to access the local storage with
  - `VecGetArray(Vec, double *[])`
- You must return the array to PETSc when you finish
  - `VecRestoreArray(Vec, double *[])`
- Allows PETSc to handle data structure conversions
  - Commonly, these routines are inexpensive and do not involve a copy



# VecGetArray in C

```
Vec          v;  
PetscScalar  *array;  
PetscInt     n, i;  
PetscErrorCode ierr;  
  
VecGetArray(v, &array);  
VecGetLocalSize(v, &n);  
PetscSynchronizedPrintf(PETSC_COMM_WORLD,  
    "First element of local array is %f\n", array[0]);  
PetscSynchronizedFlush(PETSC_COMM_WORLD);  
for(i = 0; i < n; ++i) {  
    array[i] += (PetscScalar) rank;  
}  
VecRestoreArray(v, &array);
```

# VecGetArray in F77

```
#include "finclude/petsc.h"
```

```
Vec          v;  
PetscScalar  array(1)  
PetscOffset  offset  
PetscInt     n, i  
PetscErrorCode ierr  
  
call VecGetArray(v, array, offset, ierr)  
call VecGetLocalSize(v, n, ierr)  
do i=1,n  
    array(i+offset) = array(i+offset) + rank  
end do  
call VecRestoreArray(v, array, offset, ierr)
```

# VecGetArray in F90

```
#include "finclude/petsc.h90"
```

```
Vec                v;  
PetscScalar        pointer :: array(:)  
PetscInt           n, i  
PetscErrorCode ierr  
  
call VecGetArrayF90(v, array, ierr)  
call VecGetLocalSize(v, n, ierr)  
do i=1,n  
    array(i) = array(i) + rank  
end do  
call VecRestoreArrayF90(v, array, ierr)
```

# Outline

2

## PETSc Integration

- Initial Operations
- Vector Algebra
- **Matrix Algebra**
- Algebraic Solvers
- More Abstractions

# Matrix Algebra

## What are PETSc matrices?

- Fundamental objects for storing stiffness matrices and Jacobians
- Each process locally owns a contiguous set of rows
- Supports many data types
  - AIJ, Block AIJ, Symmetric AIJ, Block Diagonal, etc.
- Supports structures for many packages
  - MUMPS, Spooles, SuperLU, UMFPack, DSCPack

# How do I create matrices?

- `MatCreate(MPI_Comm, Mat *)`
- `MatSetSizes(Mat, int m, int n, int M, int N)`
- `MatSetType(Mat, MatType typeName)`
- `MatSetFromOptions(Mat)`
  - Can set the type at runtime
- `MatSeqAIJPreallocation(Mat, PetscInt nz, const PetscInt nnz[])`
- `MatMPIAIJPreallocation(Mat, dnz, dnz[], onz, onz[])`
- `MatSetValues(Mat, m, rows[], n, cols[], values[], InsertMode)`
  - **MUST** be used, but does automatic communication

# Matrix Polymorphism

The PETSc `Mat` has a single user interface,

- Matrix assembly
  - `MatSetValues()`
- Matrix-vector multiplication
  - `MatMult()`
- Matrix viewing
  - `MatView()`

but multiple underlying implementations.

- AIJ, Block AIJ, Symmetric Block AIJ,
- Dense
- Matrix-Free
- etc.

A matrix is defined by its **interface**, not by its **data structure**.

# Matrix Assembly

- A three step process

- Each process sets or adds values
- Begin communication to send values to the correct process
- Complete the communication

- `MatSetValues(Mat m, m, rows[], n, cols[], values[], mode)`

- `mode` is either `INSERT_VALUES` or `ADD_VALUES`
- Logically dense block of values

- Two phase assembly allows overlap of communication and computation

- `MatAssemblyBegin(Mat m, type)`
- `MatAssemblyEnd(Mat m, type)`
- `type` is either `MAT_FLUSH_ASSEMBLY` or `MAT_FINAL_ASSEMBLY`



# One Way to Set the Elements of a Matrix

## Simple 3-point stencil for 1D Laplacian

```
v[0] = -1.0; v[1] = 2.0; v[2] = -1.0;
if (rank == 0) {
    for(row = 0; row < N; row++) {
        cols[0] = row-1; cols[1] = row; cols[2] = row+1;
        if (row == 0) {
            MatSetValues(A,1,&row,2,&cols[1],&v[1],INSERT_VALUES);
        } else if (row == N-1) {
            MatSetValues(A,1,&row,2,cols,v,INSERT_VALUES);
        } else {
            MatSetValues(A,1,&row,3,cols,v,INSERT_VALUES);
        }
    }
}

MatAssemblyBegin(A, MAT_FINAL_ASSEMBLY);
MatAssemblyEnd(A, MAT_FINAL_ASSEMBLY);
```

# A Better Way to Set the Elements of a Matrix

Simple 3-point stencil for 1D Laplacian

```
v[0] = -1.0; v[1] = 2.0; v[2] = -1.0;
MatGetOwnershipRange(A, &start, &end);
for(row = start; row < end; row++) {
    cols[0] = row-1; cols[1] = row; cols[2] = row+1;
    if (row == 0) {
        MatSetValues(A, 1, &row, 2, &cols[1], &v[1], INSERT_VALUES);
    } else if (row == N-1) {
        MatSetValues(A, 1, &row, 2, cols, v, INSERT_VALUES);
    } else {
        MatSetValues(A, 1, &row, 3, cols, v, INSERT_VALUES);
    }
}
MatAssemblyBegin(A, MAT_FINAL_ASSEMBLY);
MatAssemblyEnd(A, MAT_FINAL_ASSEMBLY);
```

# Why Are PETSc Matrices That Way?

- No one data structure is appropriate for all problems
  - Blocked and diagonal formats provide significant performance benefits
  - PETSc has many formats and makes it easy to add new data structures
- Assembly is difficult enough without worrying about partitioning
  - PETSc provides parallel assembly routines
  - Achieving high performance still requires making most operations local
  - However, programs can be incrementally developed.
  - `MatPartitioning` and `MatOrdering` can help
- Matrix decomposition in contiguous chunks is simple
  - Makes interoperation with other codes easier
  - For other ordering, PETSc provides “Application Orderings” (AO)

# Outline

2

## PETSc Integration

- Initial Operations
- Vector Algebra
- Matrix Algebra
- **Algebraic Solvers**
- More Abstractions

# Experimentation is Essential!

Proof is not currently enough to examine solvers

- N. M. Nachtigal, S. C. Reddy, and L. N. Trefethen, *How fast are nonsymmetric matrix iterations?*, SIAM J. Matrix Anal. Appl., **13**, pp.778–795, 1992.
- Anne Greenbaum, Vlastimil Ptak, and Zdenek Strakos, *Any Nonincreasing Convergence Curve is Possible for GMRES*, SIAM J. Matrix Anal. Appl., **17** (3), pp.465–469, 1996.

# Solver Types

- **Explicit:**
  - Field variables are updated using local neighbor information
- **Semi-implicit:**
  - Some subsets of variables are updated with global solves
  - Others with direct local updates
- **Implicit:**
  - Most or all variables are updated in a single global solve

# Linear Solvers

## Krylov Methods

- Using PETSc linear algebra, just add:

- `KSPSetOperators(KSP ksp, Mat A, Mat M, MatStructure flag)`
- `KSPSolve(KSP ksp, Vec b, Vec x)`

- Can access subobjects

- `KSPGetPC(KSP ksp, PC *pc)`

- Preconditioners must obey PETSc interface

- Basically just the KSP interface

- Can change solver dynamically from the command line

- `-ksp_type bicgstab`

# Nonlinear Solvers

## Newton and Picard Methods

- Using PETSc linear algebra, just add:

- `SNESSetFunction(SNES snes, Vec r, residualFunc, void *ctx)`
- `SNESSetJacobian(SNES snes, Mat A, Mat M, jacFunc, void *ctx)`
- `SNESolve(SNES snes, Vec b, Vec x)`

- Can access subobjects

- `SNESGetKSP(SNES snes, KSP *ksp)`

- Can customize subobjects from the cmd line

- Set the subdomain preconditioner to ILU with `-sub_pc_type ilu`



# Basic Solver Usage

We will illustrate basic solver usage with `SNES`.

- Use `SNESSetFromOptions()` so that everything is set dynamically
  - Use `-snes_type` to set the type or take the default
- Override the tolerances
  - Use `-snes_rtol` and `-snes_atol`
- View the solver to make sure you have the one you expect
  - Use `-snes_view`
- For debugging, monitor the residual decrease
  - Use `-snes_monitor`
  - Use `-ksp_monitor` to see the underlying linear solver

# 3rd Party Solvers in PETSc

## Complete table of solvers

### 1 Sequential LU

- ILUDT (SPARSEKIT2, Yousef Saad, U of MN)
- EUCLID & PILUT (Hypre, David Hysom, LLNL)
- ESSL (IBM)
- SuperLU (Jim Demmel and Sherry Li, LBNL)
- Matlab
- UMFPACK (Tim Davis, U. of Florida)
- LUSOL (MINOS, Michael Saunders, Stanford)

### 2 Parallel LU

- MUMPS (Patrick Amestoy, IRIT)
- SPOOLES (Cleve Ashcroft, Boeing)
- SuperLU\_Dist (Jim Demmel and Sherry Li, LBNL)

### 3 Parallel Cholesky

- DSCPACK (Padma Raghavan, Penn. State)

### 4 XYTLlib - parallel direct solver (Paul Fischer and Henry Tufo, ANL)

# 3rd Party Preconditioners in PETSc

## Complete table of solvers

- 1 Parallel ICC
  - BlockSolve95 (Mark Jones and Paul Plassman, ANL)
- 2 Parallel ILU
  - BlockSolve95 (Mark Jones and Paul Plassman, ANL)
- 3 Parallel Sparse Approximate Inverse
  - Parasails (Hypre, Edmund Chow, LLNL)
  - SPAI 3.0 (Marcus Grote and Barnard, NYU)
- 4 Sequential Algebraic Multigrid
  - RAMG (John Ruge and Klaus Steuben, GMD)
  - SAMG (Klaus Steuben, GMD)
- 5 Parallel Algebraic Multigrid
  - Prometheus (Mark Adams, PPPL)
  - BoomerAMG (Hypre, LLNL)
  - ML (Trilinos, Ray Tuminaro and Jonathan Hu, SNL)

# Outline

- 2 PETSc Integration
  - Initial Operations
  - Vector Algebra
  - Matrix Algebra
  - Algebraic Solvers
  - **More Abstractions**

# Higher Level Abstractions

The PETSc `DA` class is a topology and discretization interface.

- Structured grid interface
  - Fixed simple topology
- Supports stencils, communication, reordering
  - Limited idea of operators
- Nice for simple finite differences

The PETSc `Mesh` class is a topology interface.

- Unstructured grid interface
  - Arbitrary topology and element shape
- Supports partitioning, distribution, and global orders

# Higher Level Abstractions

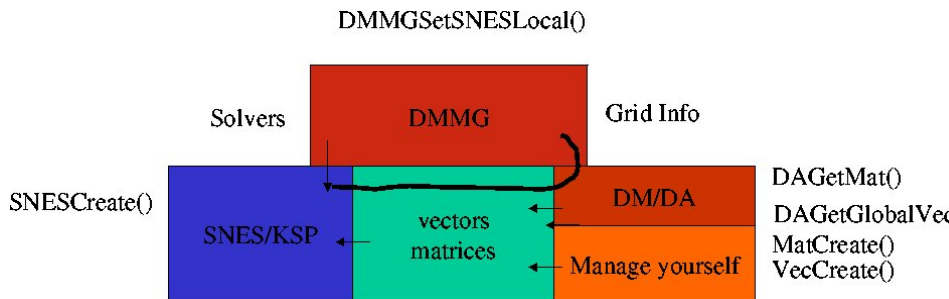
The PETSc `DM` class is a hierarchy interface.

- Supports multigrid
  - DMMG combines it with the MG preconditioner
- Abstracts the logic of multilevel methods

The PETSc `Section` class is a function interface.

- Functions over unstructured grids
  - Arbitrary layout of degrees of freedom
- Support distribution and assembly

# 3 Ways To Use PETSc



- User manages all topology (just use Vec and Mat)
- PETSc manages single topology (use DM)
- PETSc manages a hierarchy (use DMMG)

# Outline

1 Getting Started with PETSc

2 PETSc Integration

3 Common PETSc Usage

- Principles and Design
- Debugging PETSc
- Profiling PETSc
- Serial Performance
- Modeling Code

4 Advanced PETSc

5 Future Plans

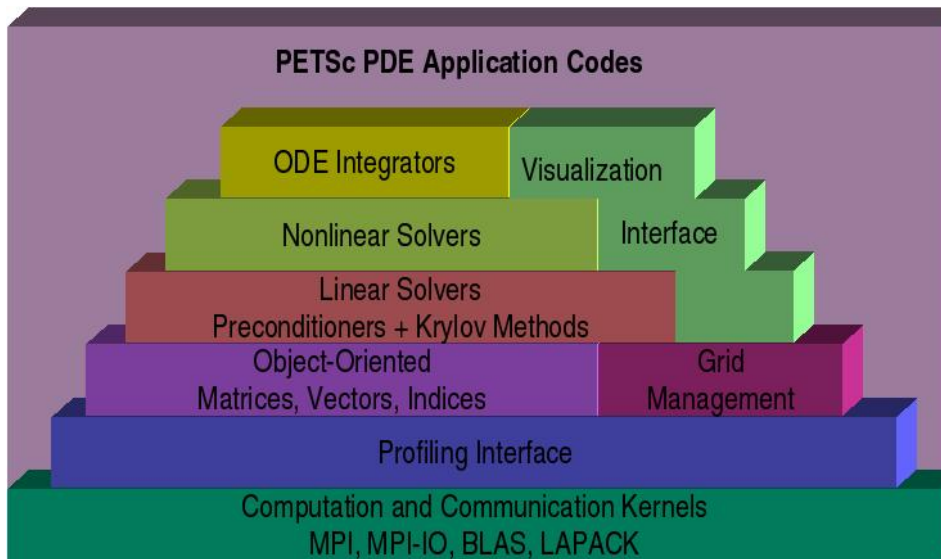


# Outline

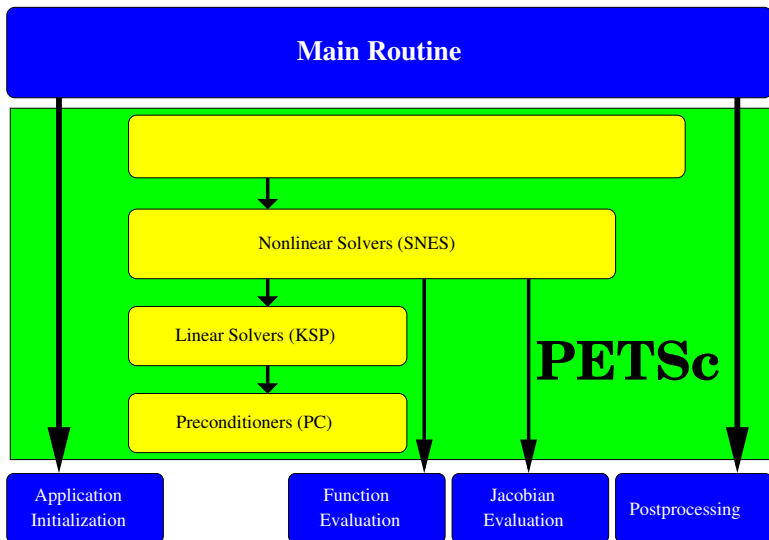
## 3 Common PETSc Usage

- Principles and Design
- Debugging PETSc
- Profiling PETSc
- Serial Performance
- Modeling Code

# PETSc Structure



# Flow Control for a PETSc Application



# Levels of Abstraction

## In Mathematical Software

- Application-specific interface
  - Programmer manipulates objects associated with the application
- High-level mathematics interface
  - Programmer manipulates mathematical objects
    - Weak forms, boundary conditions, meshes
- Algorithmic and discrete mathematics interface
  - Programmer manipulates mathematical objects
    - Sparse matrices, nonlinear equations
  - Programmer manipulates algorithmic objects
    - Solvers
- Low-level computational kernels
  - BLAS-type operations, FFT

# Object-Oriented Design

- Design based on **operations** you perform,
  - rather than the data in the object
- Example: A vector is
  - **not** a 1d array of numbers
  - **an object** allowing addition and scalar multiplication
- The efficient use of the computer is an added difficulty
  - which often leads to code generation

# What is not in PETSc?

- ~~Unstructured mesh generation and manipulation~~
  - In 3.0, we have `Mesh` objects
- Discretizations
  - `DealII`
  - In 3.0, we have an interface to `FIAT`
- Higher level representations of PDEs
  - `FEniCS` (`FFC/Syfi`) and `Sundance`
- Load balancing
  - Interface to `Zoltan`
- Sophisticated visualization capabilities
  - Interface to `MayaVi2` through VTK
- Eigenvalues
  - `SLEPc` and `SIP`
- Optimization and sensitivity
  - `TAO` and `Veltisto`

# Outline

## 3 Common PETSc Usage

- Principles and Design
- **Debugging PETSc**
- Profiling PETSc
- Serial Performance
- Modeling Code

# Correctness Debugging

- Automatic generation of tracebacks
- Detecting memory corruption and leaks
- Optional user-defined error handlers



# Interacting with the Debugger

- Launch the debugger
  - `-start_in_debugger [gdb,dbx,nxterm]`
  - `-on_error_attach_debugger [gdb,dbx,nxterm]`
- Attach the debugger only to some parallel processes
  - `-debugger_nodes 0,1`
- Set the display (often necessary on a cluster)
  - `-display khan.mcs.anl.gov:0.0`

# Debugging Tips

- Put a breakpoint in `PetscError()` to catch errors as they occur
- PETSc tracks memory overwrites at both ends of arrays
  - The `CHKMEMQ` macro causes a check of all allocated memory
  - Track memory overwrites by bracketing them with `CHKMEMQ`
- PETSc checks for leaked memory
  - Use `PetscMalloc()` and `PetscFree()` for all allocation
  - Print unfreed memory on `PetscFinalize()` with `-malloc_dump`
- Simply the best tool today is **valgrind**
  - It checks memory access, cache performance, memory usage, etc.
  - <http://www.valgrind.org>
  - Need `-trace-children=yes` when running under MPI

# Exercise 7

Use the debugger to find a SEGV  
Locate a memory overwrite using CHKMEMQ.

- Get the example

- `hg clone -r1`

`http://petsc.cs.iit.edu/petsc/SimpleTutorial`

- Build the example `make`

- Run it and watch the fireworks

- `mpiexec -n 2 ./bin/ex5 -use_coords`

- Run it under the debugger and correct the error

- `mpiexec -n 2 ./bin/ex5 -use_coords`  
`-start_in_debugger -display :0.0`
  - `hg update -r2`

- Build it and run again smoothly

# Outline

## 3 Common PETSc Usage

- Principles and Design
- Debugging PETSc
- **Profiling PETSc**
- Serial Performance
- Modeling Code

# Performance Debugging

- PETSc has integrated profiling
  - Option `-log_summary` prints a report on `PetscFinalize()`
- PETSc allows user-defined events
  - Events report time, calls, flops, communication, etc.
  - Memory usage is tracked by object
- Profiling is separated into stages
  - Event statistics are aggregated by stage

# Using Stages and Events

- Use `PetscLogStageRegister()` to create a new stage
  - Stages are identifier by an integer handle
- Use `PetscLogStagePush/Pop()` to manage stages
  - Stages may be nested, but will not aggregate in a nested fashion
- Use `PetscLogEventRegister()` to create a new stage
  - Events also have an associated class
- Use `PetscLogEventBegin/End()` to manage events
  - Events may also be nested and will aggregate in a nested fashion
  - Can use `PetscLogFlops()` to log user flops

# Adding A Logging Stage

```
int stageNum;  
  
PetscLogStageRegister(&stageNum, "name");  
PetscLogStagePush(stageNum);  
  
/* Code to Monitor */  
  
PetscLogStagePop();
```

# Adding A Logging Event

```
static int USER_EVENT;  
  
PetscLogEventRegister(&USER_EVENT, "name", CLS_ID);  
PetscLogEventBegin(USER_EVENT, 0, 0, 0, 0);  
  
/* Code to Monitor */  
  
PetscLogFlops(user_event_flops);  
PetscLogEventEnd(USER_EVENT, 0, 0, 0, 0);
```



# Adding A Logging Class

```
static int CLASS_ID;
```

```
PetscLogClassRegister(&CLASS_ID, "name");
```

- Class ID identifies a class uniquely
- Must initialize before creating any objects of this type

# Matrix Memory Preallocation

- PETSc sparse matrices are dynamic data structures
  - can add additional nonzeros freely
- Dynamically adding many nonzeros
  - requires additional memory allocations
  - requires copies
  - can kill performance
- Memory preallocation provides
  - the freedom of dynamic data structures
  - good performance
- Easiest solution is to replicate the assembly code
  - Remove computation, but preserve the indexing code
  - Store set of columns for each row
- Call preallocation routines for all datatypes
  - `MatSeqAIJSetPreallocation()`
  - `MatMPIAIJSetPreallocation()`
  - Only the relevant data will be used

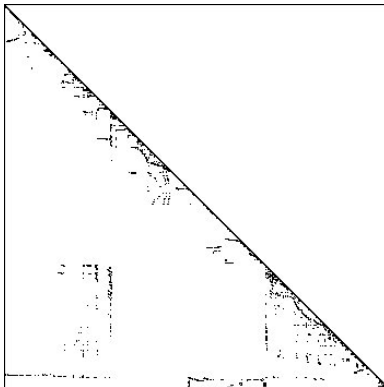
# Matrix Memory Preallocation

## Sequential Sparse Matrices

```
MatSeqAIJPreallocation(Mat A, int nz, int nnz[])
```

**nz**: expected number of nonzeros in any row

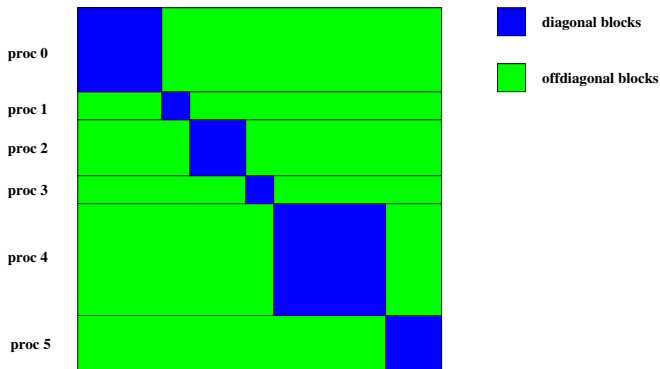
**nnz(i)**: expected number of nonzeros in row i



# Matrix Memory Preallocation

## ParallelSparseMatrix

- Each process locally owns a submatrix of contiguous global rows
- Each submatrix consists of diagonal and off-diagonal parts



- `MatGetOwnershipRange(Mat A, int *start, int *end)`

`start`: first locally owned row of global matrix

`end-1`: last locally owned row of global matrix

# Matrix Memory Preallocation

## Parallel Sparse Matrices

```
MatMPIAIJPreallocation(Mat A, int dnz, int dnnz[],  
int onz, int onnz[])
```

**dnz**: expected number of nonzeros in any row in the diagonal block

**dnnz(i)**: expected number of nonzeros in row i in the diagonal block

**onz**: expected number of nonzeros in any row in the offdiagonal portion

**onnz(i)**: expected number of nonzeros in row i in the offdiagonal portion

# Matrix Memory Preallocation

## Verifying Preallocation

- Use runtime option `-info`

- Output:

```
[proc #] Matrix size:  %d X %d; storage space:  
%d unneeded, %d used
```

```
[proc #] Number of mallocs during MatSetValues( )  
is %d
```

```
[merlin] mpirun ex2 -log_info  
[0]MatAssemblyEnd_SeqAIJ:Matrix size: 56 X 56; storage space:  
[0] 310 unneeded, 250 used  
[0]MatAssemblyEnd_SeqAIJ:Number of mallocs during MatSetValues() is 0  
[0]MatAssemblyEnd_SeqAIJ:Most nonzeros in any row is 5  
[0]Mat_AIJ_CheckInode: Found 56 nodes out of 56 rows. Not using Inode routine  
[0]Mat_AIJ_CheckInode: Found 56 nodes out of 56 rows. Not using Inode routine  
Norm of error 0.000156044 iterations 6  
[0]PetscFinalize:PETSc successfully ended!
```

# Exercise 8

Return to Exercise 7 and add more profiling.

- Update to the next revision
  - `hg update -r3`
- Build, run, and look at the profiling report
  - `make ex5`
  - `./bin/ex5 -use_coords -log_summary`
- Add a new stage for setup
- Add a new event for `FormInitialGuess()` and log the flops
- Build it again and look at the profiling report

# Outline

## 3 Common PETSc Usage

- Principles and Design
- Debugging PETSc
- Profiling PETSc
- **Serial Performance**
- Modeling Code



# STREAM Benchmark

Simple benchmark program measuring **sustainable** memory bandwidth

- Prototypical operation is Triad (WAXPY):  $\mathbf{w} = \mathbf{y} + \alpha \mathbf{x}$
- Measures the memory bandwidth bottleneck (much below peak)
- Datasets outstrip cache

Machine	Peak (MF/s)	Triad (MB/s)	MF/MW	Eq. MF/s
Matt's Laptop	1700	1122.4	12.1	93.5 (5.5%)
Intel Core2 Quad	38400	5312.0	57.8	442.7 (1.2%)
Tesla 1060C	984000	102000.0*	77.2	8500.0 (0.8%)

**Table:** Bandwidth limited machine performance

<http://www.cs.virginia.edu/stream>

# Analysis of Sparse Matvec (SpMV)

## Assumptions

- No cache misses
- No waits on memory references

## Notation

$m$  Number of matrix rows

$nz$  Number of nonzero matrix elements

$V$  Number of vectors to multiply

We can look at bandwidth needed for peak performance

$$\left(8 + \frac{2}{V}\right) \frac{m}{nz} + \frac{6}{V} \text{ byte/flop} \quad (1)$$

or achievable performance given a bandwidth  $BW$

$$\frac{Vnz}{(8V + 2)m + 6nz} BW \text{ Mflop/s} \quad (2)$$

Towards Realistic Performance Bounds for Implicit CFD Codes, Gropp, Kaushik, Keyes, and Smith.

# Improving Serial Performance

For a single matvec with 3D FD Poisson, Matt's laptop can achieve at most

$$\frac{1}{(8 + 2)^{\frac{1}{7}} + 6} \text{ bytes/flop} (1122.4 \text{ MB/s}) = \mathbf{151} \text{ MFlops/s}, \quad (3)$$

which is a dismal **8.8%** of peak.

Can improve performance by

- Blocking
- Multiple vectors

but operation issue limitations take over.

# Improving Serial Performance

For a single matvec with 3D FD Poisson, Matt's laptop can achieve at most

$$\frac{1}{(8+2)^{\frac{1}{7}}+6} \text{ bytes/flop}(1122.4 \text{ MB/s}) = \textcolor{red}{151} \text{ MFlops/s}, \quad (3)$$

which is a dismal 8.8% of peak.

Better approaches:

- Unassembled operator application (Spectral elements, FMM)
  - $N$  data,  $N^2$  computation
- Nonlinear evaluation (Picard, FAS, Exact Polynomial Solvers)
  - $N$  data,  $N^k$  computation

# Performance Tradeoffs

We must balance storage, bandwidth, and cycles

- Assembled Operator Action
  - Trades cycles and storage for bandwidth in application
- Unassembled Operator Action
  - Trades bandwidth and storage for cycles in application
  - For high orders, storage is impossible
  - Can make use of FErari decomposition to save calculation
  - Could store element matrices to save cycles
- Partial assembly gives even finer control over tradeoffs
  - Also allows introduction of parallel costs (load balance, ...)

# Outline

## 3 Common PETSc Usage

- Principles and Design
- Debugging PETSc
- Profiling PETSc
- Serial Performance
- **Modeling Code**

# Importance of Computational Modeling

Without a model,  
performance measurements are meaningless!

Before a code is written, we should have a model of

- computation
- memory usage
- communication
- bandwidth
- achievable concurrency

This allows us to

- **verify** the implementation
- **predict** scaling behavior

# Outline

- 1 Getting Started with PETSc
- 2 PETSc Integration
- 3 Common PETSc Usage
- 4 Advanced PETSc**
  - SNES
  - DA
- 5 Future Plans
- 6 Conclusions



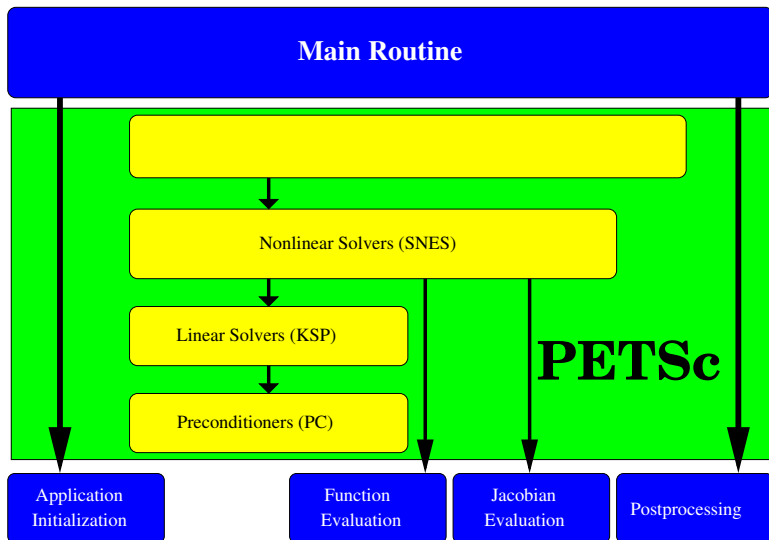
# Outline

## 4 Advanced PETSc

- SNES

- DA

# Flow Control for a PETSc Application



# SNES Paradigm

The SNES interface is based upon callback functions

- `FormFunction()`, **set by** `SNESSetFunction()`
- `FormJacobian()`, **set by** `SNESSetJacobian()`

When PETSc needs to evaluate the nonlinear residual  $F(x)$ ,

- Solver calls the **user's** function
- User function gets application state through the `ctx` variable
  - PETSc never sees application data

# Topology Abstractions

- DA

- Abstracts Cartesian grids in any dimension
- Supports stencils, communication, reordering
- Nice for simple finite differences

- Mesh

- Abstracts general topology in any dimension
- Also supports partitioning, distribution, and global orders
- Allows arbitrary element shapes and discretizations

# Assembly Abstractions

- DM
  - Abstracts the logic of multilevel (multiphysics) methods
  - Manages allocation and assembly of local and global structures
  - Interfaces to `DMMG` solver
- Section
  - Abstracts functions over a topology
  - Manages allocation and assembly of local and global structures
  - Will merge with `DM` somehow

# SNES Function

The user provided function which calculates the nonlinear residual has signature

```
PetscErrorCode (*func)(SNES snes, Vec x, Vec r, void *ctx)
```

**x**: The current solution

**r**: The residual

**ctx**: The user context passed to `SNESSetFunction()`

- Use this to pass application information, e.g. physical constants

# SNES Jacobian

The user provided function which calculates the Jacobian has signature

```
(*func) (SNES snes, Vec x, Mat *J, Mat *M, MatStructure *flag, void *ctx)
```

`x`: The current solution

`J`: The Jacobian

`M`: The Jacobian preconditioning matrix (possibly J itself)

`ctx`: The user context passed to `SNESSetJacobian()`

- Use this to pass application information, e.g. physical constants

- Possible `MatStructure` values are:

- `SAME_NONZERO_PATTERN`
- `DIFFERENT_NONZERO_PATTERN`

Alternatively, you can use

- a builtin sparse finite difference approximation
- automatic differentiation (ADIC/ADIFOR)

# SNES Variants

- Line search strategies
- Trust region approaches
- Pseudo-transient continuation
- Matrix-free variants



# Finite Difference Jacobians

PETSc can compute and explicitly store a Jacobian via 1st-order FD

- Dense
  - Activated by `-snes_fd`
  - Computed by `SNESDefaultComputeJacobian()`
- Sparse via colorings
  - Coloring is created by `MatFDColoringCreate()`
  - Computed by `SNESDefaultComputeJacobianColor()`

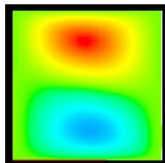
Can also use Matrix-free Newton-Krylov via 1st-order FD

- Activated by `-snes_mf` without preconditioning
- Activated by `-snes_mf_operator` with user-defined preconditioning
  - Uses preconditioning matrix from `SNESSetJacobian()`

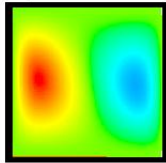
# SNES Example

## Driven Cavity

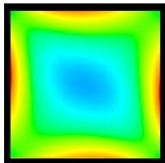
### Solution Components



velocity:  $u$



velocity:  $v$



vorticity:



temperature:  $T$

- Velocity-vorticity formulation
- Flow driven by lid and/or bouyancy
- Logically regular grid
  - Parallelized with DA
- Finite difference discretization
- Authored by David Keyes

`$PETCS_DIR/src/snes/examples/tutorials/ex19.c`

# Driven Cavity Application Context

```
typedef struct {  
    /*----- basic application data -----*/  
    double    lid_velocity;  
    double    prandtl, grashof;  
    int        mx, my;  
    int        mc;  
    PetscTruth draw_contours;  
    /*----- parallel data -----*/  
    MPI_Comm   comm;  
    DA          da;  
    /* Local ghosted solution and residual */  
    Vec         localX, localF;  
} AppCtx;
```

\$PETCS\_DIR/src/snes/examples/tutorials/ex19.c

# Driven Cavity Residual Evaluation

```
Residual(SNES snes, Vec X, Vec F, void *ptr) {  
    AppCtx          *user = (AppCtx *) ptr;
```

```
    /* local starting and ending grid points */
```

```
    int              istart, iend, jstart, jend;
```

```
    PetscScalar      *f; /* local vector data */
```

```
    PetscReal        grashof = user->grashof;
```

```
    PetscReal        prandtl = user->prandtl;
```

```
    PetscErrorCode    ierr;
```

```
    /* Code to communicate nonlocal ghost point data */
```

```
    VecGetArray(F, &f);
```

```
    /* Code to compute local function components */
```

```
    VecRestoreArray(F, &f);
```

```
    return 0;
```

```
}
```

# Better Driven Cavity Residual Evaluation

```
ResLocal(DALocalInfo *info, Field **x, Field **f, void *c)
/* Handle boundaries */
/* Compute over the interior points */
for(j = info->ys; j < info->xs+info->xm; j++) {
    for(i = info->xs; i < info->ys+info->ym; i++) {
        /* U velocity */
        u    = x[j][i].u;
        uxx  = (2.0*u - x[j][i-1].u - x[j][i+1].u)*hydhx;
        uyy  = (2.0*u - x[j-1][i].u - x[j+1][i].u)*hxdhy;
        upw  = 0.5*(x[j+1][i].omega-x[j-1][i].omega)*hx
        f[j][i].u = uxx + uyy - upw;
        /* V velocity, Omega, Temperature */
    }
}
}
```

\$PETCS\_DIR/src/snes/examples/tutorials/ex19.c

# Outline

## 4 Advanced PETSc

- SNES

- DA

# What is a DA?

DA is a topology interface handling parallel data layout on structured grids

- Handles local and global indices
  - `DAGetGlobalIndices()` and `DAGetAO()`
- Provides local and global vectors
  - `DAGetGlobalVector()` and `DAGetLocalVector()`
- Handles ghost values coherence
  - `DAGetGlobalToLocal()` and `DAGetLocalToGlobal()`

# DA Paradigm

The DA interface is based upon *local* callback functions

- `FormFunctionLocal()`, **set by** `DASetLocalFunction()`
- `FormJacobianLocal()`, **set by** `DASetLocalJacobian()`

When PETSc needs to evaluate the nonlinear residual  $F(x)$ ,

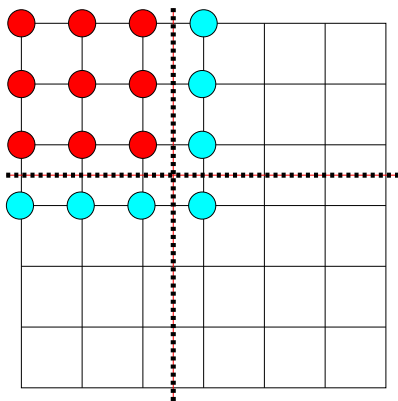
- Each process evaluates the local residual
- PETSc assembles the global residual automatically
  - Uses `DALocalToGlobal()` method



# Ghost Values

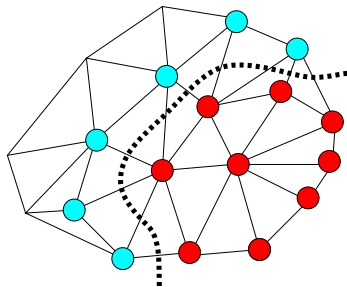
To evaluate a local function  $f(x)$ , each process requires

- its local portion of the vector  $x$
- its **ghost values**, bordering portions of  $x$  owned by neighboring processes



● Local Node

● Ghost Node



# DA Global Numberings

Proc 2			Proc 3	
25	26	27	28	29
20	21	22	23	24
15	16	17	18	19
10	11	12	13	14
5	6	7	8	9
0	1	2	3	4
Proc 0			Proc 1	

Natural numbering

Proc 2			Proc 3	
21	22	23	28	29
18	19	20	26	27
15	16	17	24	25
6	7	8	13	14
3	4	5	11	12
0	1	2	9	10
Proc 0			Proc 1	

PETSc numbering

# DA Global vs. Local Numbering

- **Global:** Each vertex has a unique id belongs on a unique process
- **Local:** Numbering includes vertices from neighboring processes
  - These are called **ghost** vertices

Proc 2			Proc 3	
X	X	X	X	X
X	X	X	X	X
12	13	14	15	X
8	9	10	11	X
4	5	6	7	X
0	1	2	3	X
Proc 0			Proc 1	

Local numbering

Proc 2			Proc 3	
21	22	23	28	29
18	19	20	26	27
15	16	17	24	25
6	7	8	13	14
3	4	5	11	12
0	1	2	9	10
Proc 0			Proc 1	

Global numbering

# DA Local Function

The user provided function which calculates the nonlinear residual in 2D has signature

```
(*lfunc) (DALocalInfo *info, PetscScalar **x, PetscScalar **r, void *ctx)
```

**info:** All layout and numbering information

**x:** The current solution

- Notice that it is a multidimensional array

**r:** The residual

**ctx:** The user context passed to `DASetLocalFunction()`

The local DA function is activated by calling

```
SNESSetFunction(snes, r, SNESDAFormFunction, ctx)
```

# Bratu Residual Evaluation

$$\Delta u + \lambda e^u = 0$$

```

ResLocal(DALocalInfo *info, Field **x, Field **f, void *c
{
    /* Not Shown: Handle boundaries */
    /* Compute over the interior points */
    for(j = info->ys; j < info->xs+info->ym; j++) {
        for(i = info->xs; i < info->ys+info->xm; i++) {
            u          = x[j][i];
            u_xx       = (2.0*u - x[j][i-1] - x[j][i+1])*hydhx;
            u_yy       = (2.0*u - x[j-1][i] - x[j+1][i])*hxdhy;
            f[j][i] = u_xx + u_yy - hx*hy*lambda*exp(u);
        }
    }
}

```

\$PETCS\_DIR/src/snes/examples/tutorials/ex5.c

# DA Local Jacobian

The user provided function which calculates the Jacobian in 2D has signature

```
(*lfunc) (DALocalInfo *info, PetscScalar **x, Mat J, void *ctx)
```

**info:** All layout and numbering information

**x:** The current solution

**J:** The Jacobian

**ctx:** The user context passed to `DASetLocalJacobian()`

The local DA function is activated by calling

```
SNESSetJacobian(snes, J, J, SNESDAComputeJacobian, ctx)
```

# Bratu Jacobian Evaluation

```
JacLocal(DALocalInfo *info, PetscScalar **x, Mat jac, void  
for(j = info->ys; j < info->ys + info->ym; j++) {  
    for(i = info->xs; i < info->xs + info->xm; i++) {  
        row.j = j; row.i = i;  
        if (i == 0 || j == 0 || i == mx-1 || j == my-1) {  
            v[0] = 1.0;  
            MatSetValuesStencil(jac, 1, &row, 1, &row, v, INSERT_VALUES);  
        } else {  
            v[0] = -(hx/hy); col[0].j = j-1; col[0].i = i;  
            v[1] = -(hy/hx); col[1].j = j; col[1].i = i-1;  
            v[2] = 2.0*(hy/hx+hx/hy)  
                - hx*hy*lambda*PetscExpScalar(x[j][i]);  
            v[3] = -(hy/hx); col[3].j = j; col[3].i = i+1;  
            v[4] = -(hx/hy); col[4].j = j+1; col[4].i = i;  
            MatSetValuesStencil(jac, 1, &row, 5, col, v, INSERT_VALUES);  
        }  
    }  
}
```

# A DA is more than a Mesh

A DA contains **topology**, **geometry**, and an implicit Q1 **discretization**.

It is used as a template to create

- Vectors (functions)
- Matrices (linear operators)



# DA Vectors

- The DA object contains only layout (topology) information
  - All field data is contained in PETSc `Vecs`
- Global vectors are parallel
  - Each process stores a unique local portion
  - `DACreateGlobalVector(DA da, Vec *gvec)`
- Local vectors are sequential (and usually temporary)
  - Each process stores its local portion plus ghost values
  - `DACreateLocalVector(DA da, Vec *lvec)`
  - includes ghost values!

# Updating Ghosts

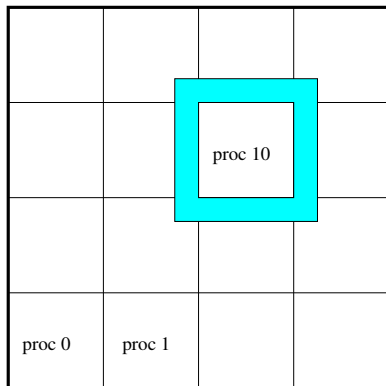
Two-step process enables overlapping computation and communication

- `DAGlobalToLocalBegin(da, gvec, mode, lvec)`
  - `gvec` provides the data
  - `mode` is either `INSERT_VALUES` or `ADD_VALUES`
  - `lvec` holds the local and ghost values
- `DAGlobalToLocalEnd(da, gvec, mode, lvec)`
  - Finishes the communication

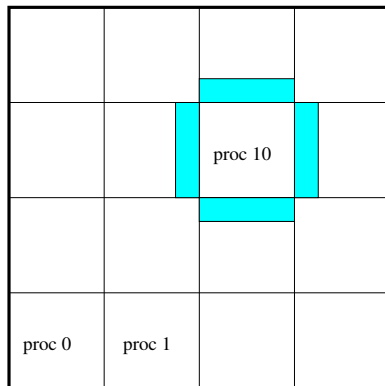
The process can be reversed with `DALocalToGlobal()`.

# DA Stencils

Both the **box** stencil and **star** stencil are available.



Box Stencil



Star Stencil

# Setting Values on Regular Grids

## PETSc provides

```
MatSetValuesStencil(Mat A, m, MatStencil idxm[], n, MatStencil idxn[],  
                    PetscScalar values[], InsertMode mode)
```

- Each row or column is actually a `MatStencil`
  - This specifies grid coordinates and a component if necessary
  - Can imagine for unstructured grids, they are *vertices*
- The values are the same logically dense block in row/col

# Creating a DA

```
DACreate2d(comm, wrap, type, M, N, m, n, dof, s, lm[], ln[], DA *da)
```

**wrap:** Specifies periodicity

- `DA_NONPERIODIC`, `DA_XPERIODIC`, `DA_YPERIODIC`, or `DA_XYPERIODIC`

**type:** Specifies stencil

- `DA_STENCIL_BOX` or `DA_STENCIL_STAR`

**M/N:** Number of grid points in x/y-direction

**m/n:** Number of processes in x/y-direction

**dof:** Degrees of freedom per node

**s:** The stencil width

**lm/n:** Alternative array of local sizes

- Use `PETSC_NULL` for the default

# Outline

1 Getting Started with PETSc

2 PETSc Integration

3 Common PETSc Usage

4 Advanced PETSc

5 **Future Plans**

- PCFieldSplit
- PETSc-GPU
- DMMG
- Mesh
- FEniCS Tools
- PetFMM

# Things To Check Out

- PCFieldSplit for multiphysics
- GPU acceleration for KSP
- DMMG for multilevel solvers
- Sieve for topology automation
- DealII and FEniCS for FEM automation
- PetFMM for particle methods

# Outline

5

## Future Plans

- PCFieldSplit
- PETSc-GPU
- DMMG
- Mesh
- FEniCS Tools
- PetFMM



# MultiPhysics Paradigm

The PCFieldSplit interface uses the `VecScatter` objects to

- extract functions/operators corresponding to each physics
  - Local evaluation for each equation
- assemble functions/operators over all physics
  - Generalizes `LocalToGlobal()`

Notice that this works in exactly the same manner as

- multiple resolutions (MG, FMM, Wavelets)
- multiple domains (Domain Decomposition)
- multiple dimensions (ADI)

# Preconditioning

Several varieties of preconditioners can be supported:

- Block Jacobi or Block Gauss-Siedel
- Schur complement
- Block ILU (approximate coupling and Schur complement)
- Dave May's implementation of Elman-Wathen type PCs

which only require actions of individual operator blocks

Notice also that we may have any combination of

- “canned” PCs (ILU, AMG)
- PCs needing special information (MG, FMM)
- custom PCs (physics-based preconditioning, Born approximation)

since we have access to an algebraic interface

# Stokes Example

The common block preconditioners for Stokes require only options:

```
-pc_type fieldsplit  
-field_split_type multiplicative  
-fieldsplit_0_pc_type ml  
-fieldsplit_0_ksp_type preonly  
-fieldsplit_1_pc_type jacobi  
-fieldsplit_1_ksp_type preonly
```

- Also additive, symmetric\_multiplicative, and schur
- Specify DA splits using, e.g. `-pc_fieldsplit_2_fields 2,3`
- For unstructured grid use `IS` directly
- Has flexible Schur factorization and preconditioner types

FieldSplit provides the **building blocks**  
for multiphysics preconditioning.

# Stokes Example

The common block preconditioners for Stokes require only options:

```
-pc_type fieldsplit  
-field_split_type multiplicative  
-fieldsplit_0_pc_type ml  
-fieldsplit_0_ksp_type preonly  
-fieldsplit_1_pc_type jacobi  
-fieldsplit_1_ksp_type preonly
```

- Also additive, symmetric\_multiplicative, and schur
- Specify DA splits using, e.g. `-pc_fieldsplit_2_fields 2,3`
- For unstructured grid use `IS` directly
- Has flexible Schur factorization and preconditioner types

FieldSplit provides the **building blocks**  
for multiphysics preconditioning.

# Stokes Example

The common block preconditioners for Stokes require only options:

```
-pc_type fieldsplit  
-field_split_type multiplicative  
-fieldsplit_0_pc_type ml  
-fieldsplit_0_ksp_type preonly  
-fieldsplit_1_pc_type jacobi  
-fieldsplit_1_ksp_type preonly
```

- Also additive, symmetric\_multiplicative, and schur
- Specify DA splits using, e.g. `-pc_fieldsplit_2_fields 2,3`
- For unstructured grid use `IS` directly
- Has flexible Schur factorization and preconditioner types

FieldSplit provides the **building blocks** for multiphysics preconditioning.

# Outline

5

## Future Plans

- PCFieldSplit
- **PETSc-GPU**
- DMMG
- Mesh
- FEniCS Tools
- PetFMM

# PETSc-GPU

## PETSc now has support for Krylov solves on the GPU

- `-with-cuda=1 -with-cusp=1 -with-thrust=1`
  - Also possibly `-with-precision=single`
- New classes `VECCUDA` and `MATAIJCUDA`
  - Just change type on command line, `-vec_type veccuda`
- Uses **Thrust** and **Cusp** libraries from Nvidia guys
- Does not communicate vectors during solve

# Outline

5

## Future Plans

- PCFieldSplit
- PETSc-GPU
- **DMMG**
- Mesh
- FEniCS Tools
- PetFMM



# DMMG Paradigm

The DMMG interface uses the *local* DA/Mesh callback functions to

- assemble global functions/operators from local pieces
- assemble functions/operators on coarse grids

DMMG relies upon DA/Mesh (DM) to organize the

- assembly
- coarsening/refinement

while it organizes the control flow for the multilevel solve.

# DMMG Integration with SNES

- DMMG supplies global residual and Jacobian to SNES
  - User supplies local version to DMMG
  - The `Rhs_*`() and `Jac_*`() functions in the example
- Allows automatic parallelism
- Allows grid hierarchy
  - Enables multigrid once interpolation/restriction is defined
- Paradigm is developed in unstructured work
  - Solve needs scatter into contiguous global vectors (initial guess)
- Handle Neumann BC using `DMMGSetNullSpace()`

# Multigrid with DMMG

Allows multigrid with some simple command line options

- `-dmmg_nlevels`
- `-pc_mg_type, -pc_mg_cycle_type`
- `-mg_levels_1_ksp_type, -mg_levels_1_pc_type`
- `-mg_coarse_ksp_type, -mg_coarse_pc_type`
- `-dmmg_view`

Interface also works with 3rd party packages, like ML from Sandia

# Outline

5

## Future Plans

- PCFieldSplit
- PETSc-GPU
- DMMG
- **Mesh**
- FEniCS Tools
- PetFMM

# Mesh Paradigm

The Mesh interface also uses *local* callback functions

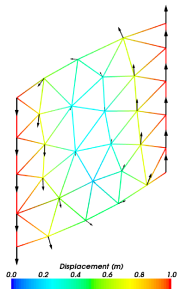
- maps between **global** Vec and **local** Vec
- Local vectors are combined into a `Section` object

When PETSc needs to evaluate the nonlinear residual  $F(x)$ ,

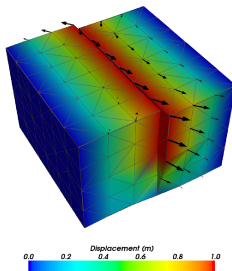
- Each process evaluates the local residual for each element
- PETSc assembles the global residual automatically
  - `SectionComplete()` generalizes `DALocalToGlobal()`

# Multiple Mesh Types

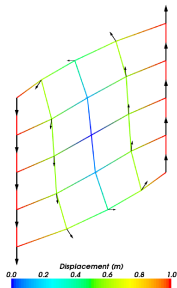
Triangular



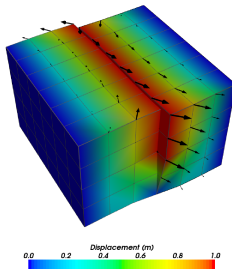
Tetrahedral



Rectangular

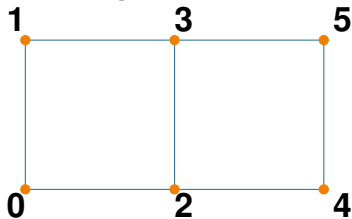


Hexahedral

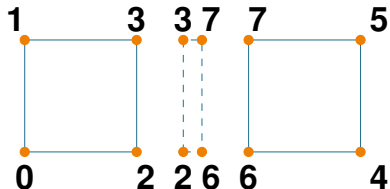
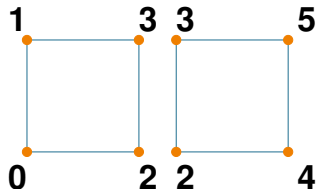
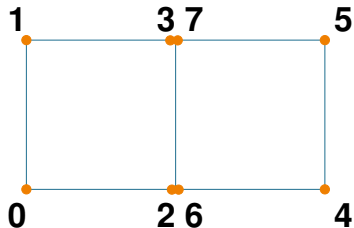


# Cohesive Cells

## Original Mesh



## Mesh with Cohesive Cell



## Exploded view of meshes

# Cohesive Cells

Cohesive cells are used to enforce slip conditions on a fault

- Demand complex mesh manipulation
  - We allow specification of only fault vertices
  - Must “sew” together on output
- Use Lagrange multipliers to enforce constraints
  - Forces illuminate physics
- Allow different fault constitutive models
  - Simplest is enforced slip
  - Now have fault constitutive models



# Outline

5

## Future Plans

- PCFieldSplit
- PETSc-GPU
- DMMG
- Mesh
- **FEniCS Tools**
- PetFMM

# FIAT

## Finite Element Integrator And Tabulator by Rob Kirby

<http://www.fenics.org/fiat>

FIAT understands

- Reference element shapes (line, triangle, tetrahedron)
- Quadrature rules
- Polynomial spaces
- Functionals over polynomials (dual spaces)
- Derivatives

User can build arbitrary elements specifying the Ciarlet triple  $(K, P, P')$

FIAT is part of the FEniCS project, as is the PETSc Sieve module

# FFC

FFC is a compiler for variational forms by Anders Logg.

Here is a mixed-form Poisson equation:

$$a((\tau, w), (\sigma, u)) = L((\tau, w)) \quad \forall (\tau, w) \in V$$

where

$$\begin{aligned} a((\tau, w), (\sigma, u)) &= \int_{\Omega} \tau \sigma - \nabla \cdot \tau u + w \nabla \cdot u \, dx \\ L((\tau, w)) &= \int_{\Omega} w f \, dx \end{aligned}$$

# FFC

## Mixed Poisson

```
shape = "triangle"
```

```
BDM1 = FiniteElement("Brezzi-Douglas-Marini", shape, 1)
```

```
DG0 = FiniteElement("Discontinuous Lagrange", shape, 0)
```

```
element = BDM1 + DG0
```

```
(tau, w) = TestFunctions(element)
```

```
(sigma, u) = TrialFunctions(element)
```

```
a = (dot(tau, sigma) - div(tau)*u + w*div(sigma))*dx
```

```
f = Function(DG0)
```

```
L = w*f*dx
```

# FFC

Here is a discontinuous Galerkin formulation of the Poisson equation:

$$a(v, u) = L(v) \quad \forall v \in V$$

where

$$\begin{aligned} a(v, u) &= \int_{\Omega} \nabla u \cdot \nabla v \, dx \\ &+ \sum_S \int_S - \langle \nabla v \rangle \cdot [[u]]_n - [[v]]_n \cdot \langle \nabla u \rangle - (\alpha/h)vu \, dS \\ &+ \int_{\partial\Omega} -\nabla v \cdot [[u]]_n - [[v]]_n \cdot \nabla u - (\gamma/h)vu \, ds \\ L(v) &= \int_{\Omega} vf \, dx \end{aligned}$$

# FFC

## DG Poisson

```
DG1 = FiniteElement("Discontinuous Lagrange", shape, 1)
v = TestFunctions(DG1)
u = TrialFunctions(DG1)
f = Function(DG1)
g = Function(DG1)
n = FacetNormal("triangle")
h = MeshSize("triangle")
a = dot(grad(v), grad(u))*dx
    - dot(avg(grad(v)), jump(u, n))*dS
    - dot(jump(v, n), avg(grad(u)))*dS
    + alpha/h*dot(jump(v, n) + jump(u, n))*dS
    - dot(grad(v), jump(u, n))*ds
    - dot(jump(v, n), grad(u))*ds
    + gamma/h*v*u*ds
L = v*f*dx + v*g*ds
```

# Outline

5

## Future Plans

- PCFieldSplit
- PETSc-GPU
- DMMG
- Mesh
- FEniCS Tools
- **PetFMM**

# FMM Applications

FMM can accelerate both integral and boundary element methods for:

- Laplace
- Stokes
- Elasticity



# FMM Applications

FMM can accelerate both integral and boundary element methods for:

- Laplace
- Stokes
- Elasticity

## Advantages

- Mesh-free
- $\mathcal{O}(N)$  time
- Distributed and multicore (GPU) parallelism
- Small memory bandwidth requirement

# PetFMM

PetFMM is an freely available implementation of the  
**F**ast **M**ultipole **M**ethod

[http://barbagroup.bu.edu/Barba\\_group/PetFMM.html](http://barbagroup.bu.edu/Barba_group/PetFMM.html)

- Leverages **PETSc**
  - Same open source license
  - Uses Sieve for parallelism
- Extensible design in C++
  - Templated over the kernel
  - Templated over traversal for evaluation
- MPI implementation
  - Novel parallel strategy for anisotropic/sparse particle distributions
  - **PetFMM—A dynamically load-balancing parallel fast multipole library**
  - 86% efficient **strong** scaling on 64 procs
- Example application using the Vortex Method for fluids
- (coming soon) GPU implementation

# Outline

- 1 Getting Started with PETSc
- 2 PETSc Integration
- 3 Common PETSc Usage
- 4 Advanced PETSc
- 5 Future Plans
- 6 Conclusions**

# Conclusions

PETSc can help you

- easily construct a code to test your ideas
  - Lots of code construction, management, and debugging tools
- scale an existing code to large or distributed machines
  - Using `FormFunctionLocal()` and scalable linear algebra
- incorporate more scalable or higher performance algorithms
  - Such as domain decomposition or multigrid
- tune your code to new architectures
  - Using profiling tools and specialized implementations

# Conclusions

PETSc can help you

- easily construct a code to test your ideas
  - Lots of code construction, management, and debugging tools
- scale an existing code to large or distributed machines
  - Using `FormFunctionLocal()` and scalable linear algebra
- incorporate more scalable or higher performance algorithms
  - Such as domain decomposition or multigrid
- tune your code to new architectures
  - Using profiling tools and specialized implementations

# Conclusions

PETSc can help you

- easily construct a code to test your ideas
  - Lots of code construction, management, and debugging tools
- scale an existing code to large or distributed machines
  - Using `FormFunctionLocal()` and scalable linear algebra
- incorporate more scalable or higher performance algorithms
  - Such as domain decomposition or multigrid
- tune your code to new architectures
  - Using profiling tools and specialized implementations

# Conclusions

PETSc can help you

- easily construct a code to test your ideas
  - Lots of code construction, management, and debugging tools
- scale an existing code to large or distributed machines
  - Using `FormFunctionLocal()` and scalable linear algebra
- incorporate more scalable or higher performance algorithms
  - Such as domain decomposition or multigrid
- tune your code to new architectures
  - Using profiling tools and specialized implementations

# Conclusions

PETSc can help you

- easily construct a code to test your ideas
  - Lots of code construction, management, and debugging tools
- scale an existing code to large or distributed machines
  - Using `FormFunctionLocal()` and scalable linear algebra
- incorporate more scalable or higher performance algorithms
  - Such as domain decomposition or multigrid
- tune your code to new architectures
  - Using profiling tools and specialized implementations